TurboQuant v1.3.0

Work Log

2026-04-01 14:00 ~ 23:00

Total 9 hours

head_dim P1-P5 + All Community Issues Resolved + TheTom Fork Analysis

AmesianX

Session Overview

- 14:00 ~ 14:30 — Work log PDF + GitHub Actions cleanup
- 14:30 ~ 15:00 — Community feedback analysis (6 issues)
- 15:00 ~ 15:30 — v1.2.0 build issues + standalone build
- 15:30 ~ 18:00 — head_dim P1-P5 Priority Cascade implementation
- 18:00 ~ 19:00 — turbo4 PPL + "////" reproduction test
- 19:00 ~ 20:00 — v1.3.0 release + community sharing
- 20:00 ~ 21:00 — Distill benchmark + scientist name test
- 21:00 ~ 22:00 — Code review + community response
- 22:00 ~ 23:00 — TheTom fork turbo4 analysis + comment correction

head_dim P1-P5 Priority Cascade Core

Before: Only read GGUF `attention.key_length` → TurboQuant silently disabled on most models

PRIORITY	SIGNAL	CONFIDENCE	NOTES
P1	`attention.key_length`	100%	GGUF official value
P2	`key_length_mla`	100%	DeepSeek V2 etc.
P3	`key_length_swa`	100%	Gemma 2/3 etc.
P4	`value_length`	95%	Cross-validation
P5	`n_embd / n_head`	70%	Wrong for MoE!

P5 Trap: Qwen3-MoE has n_embd/n_head=64 but actual head_dim=**128**. Qwen3.5-27B has 213 but actual **256**.

All Community Issues Resolved

Done

ISSUE	REPORTER	STATUS
Phi-4, DeepSeek head_dim detection failure	@fritolays	**Fixed**
turbo4-K PPL explosion (18,202)	@TheTom	**Fixed**
GLM head_dim=576, Qwen3-4B head_dim=80	@fritolays, @sztlink	**Fixed**
"////" output (Qwen3.5-27B-UD)	@modderBUG	Not reproducible
llama-bench TBQ unsupported	@sztlink	**Fixed**
Windows OpenSSL DLL dependency	@sztlink	**Fixed**

turbo4-K PPL + "////" Output Verification Resolved

TURBO4-K PPL (QWEN3-30B-A3B)

f16/f16	**6.26**	baseline
tbq4/tbq4	**6.73**	+7.5%
tbq3/tbq3	8.49	+35.6%

Normal — no explosion!

"////" OUTPUT REPRODUCTION (QWEN3.5-27B-UD)

- Tested with tbqp3/tbq3
- Normal output confirmed
- P1 correctly detected head_dim=256
- n_embd/n_head=213 → P5 was wrong

Not reproducible — environment-specific

Conclusion: Both issues had head_dim misdetection as root cause. Resolved by P1→P5 cascade.

Distill Model Benchmark New

Qwen3.5-27B-Claude-4.6-Opus-Reasoning-Distilled Q4_K_M, head_dim=256

CONFIG	PPL	VS F16	KV MEMORY	COMPRESSION
f16/f16	**6.16**	baseline	8,192 MiB	1.0x
tbqp4/tbq4	**6.20**	**+0.6%**	2,096 MiB	**3.9x**
tbqp3/tbq3	**6.23**	**+1.1%**	1,584 MiB	**5.2x**
tbq4/tbq4	7.01	+13.8%	—	—
tbq3/tbq3	7.11	+15.4%	—	—

head_dim=256 + QJL: 5.2x compression with only PPL +1.1% — nearly lossless

Scientist Name Test (6 names) Passed

tbqp3/tbq3 (5.2x compression), Qwen3.5-27B-Claude-4.6-Opus-Distilled

SCIENTIST	OUTPUT (KOREAN)	STANDARD (KOREAN)	VERDICT
Wolfgang Pauli	Bolpeugang Paulli	Bolpeugang Paulli	PASS
Niels Bohr	Nielseu Boeo	Nilseu Boeo	OK*
Erwin Schrödinger	Ereubin Syuroedingeo	Ereubin Syuroedingeo	PASS
Werner Heisenberg	Bereneo Haijenberekeu	Bereneo Haijenberekeu	PASS
Max Planck	Makseu Peullangkeu	Makseu Peullangkeu	PASS
Enrico Fermi	Eriko Pereumi	Enriko Pereumi	OK*

* Commonly used transliteration with minor difference from standard — Chinese LLM characteristic (unrelated to TurboQuant, F16 produces same output)

Code Review Fixes Quality

CRITICAL (2 ITEMS)

- Missing `head_dim=0` reset after pow2 validation
- V fallback type mismatch (f16 vs q8_0) unified

MEDIUM (1 ITEM)

- Skip entire detection logic when TBQ unused
- Remove unnecessary metadata reads + logs

After fixes: rebuild + Pauli test passed + patches/modified_files regenerated

Community Response

Communication

- **@dagbdagb** — Pointed out "6 signals" error → corrected to "5-level cascade"
- **@sztlink** — Thanks for RTX 4090 bench + requested v1.3.0 retest
- **@zekrom-vale** — turbo4 explosion was head_dim misdetection, our fork is fine
- **@Defilan** — SM 120 testing welcome + shared upstream integration plan
- **@fritolays** — Thanks for 7-architecture bench + requested Phi-4/DeepSeek retest
- **@yajo** — Explained orthogonal relationship between BitNet and TurboQuant

> **Upstream plan:** PR to llama.cpp main after stabilization. Currently managing patches+diffs in turboquant/ directory.

TheTom Fork Analysis Analysis

- **Analyzed TheTom's turboquant_plus fork code**
- **Actual root cause of turbo4-K PPL explosion:** 7 kernel bugs (SET_ROWS packing mismatch, QJL stage missing, etc.)
- **Our fork is an independent implementation** — unaffected by those bugs
- **Comment corrected:** "head_dim misdetection" → "TheTom fork's own kernel bugs"
- **Shared analysis results with TheTom** + acknowledged turbo4-resurrection paper

> **Key finding:** turbo4 PPL explosion was caused by 7 kernel bugs in TheTom's fork. Unrelated to our independent implementation.

Key Achievements Summary

7	9h	v1.3.0
Issues Resolved	Work Duration	Release Deployed

- **P1→P5 Priority Cascade** — 6-signal cross-validation
- **Community issues 6/6 resolved** — Phi-4, turbo4, GLM, ////, llama-bench, OpenSSL
- **Distill benchmark** — PPL 6.23 (+1.1%), 5.2x compression nearly lossless
- **Scientist names 6/6** — head_dim=256 QJL precision perfectly maintained
- **Code review applied** — 2 Critical + 1 Medium fixes
- **Responded to 8 community members** — shared upstream integration plan
- **TheTom fork analysis** — turbo4 PPL explosion caused by 7 kernel bugs, confirmed unrelated to our independent implementation

Thank You

TurboQuant v1.3.0

2026-04-01